

EDA_practice

January 25, 2021

1 Exploratory Data Analysis for Machine Learning

1.1 Peer Review Project 1

This project uses the publically available **Breast Cancer Wisconsin (Diagnostic) Data set** hosted on Kaggle

<https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>

Features are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image.

K. P. Bennett and O. L. Mangasarian: “Robust Linear Programming Discrimination of Two Linearly Inseparable Sets”, Optimization Methods and Software 1, 1992, 23-34.

I intend to apply the steps taught in the last 2 weeks ie;

1. Data retrieval
2. Data cleaning (removing/imputing missing values and outliers if present)
3. Checking features using plots/ visualisations
4. Feature engineering such as transformations for linear regression modelling
5. Hypothesis testing for features that are associated with Benign vs Malignant tissue

1.1.1 Module Imports

```
[1]: %pylab inline
      %config InlineBackend.figure_formats = ['retina']

      import os
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns
      sns.set()
```

Populating the interactive namespace from numpy and matplotlib

1.1.2 Importing the data

```
[2]: # Importing the data
      filepath = "data.csv"
      df = pd.read_csv(filepath)
      df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 33 columns):
```

#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	object
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64
16	smoothness_se	569 non-null	float64
17	compactness_se	569 non-null	float64
18	concavity_se	569 non-null	float64
19	concave points_se	569 non-null	float64
20	symmetry_se	569 non-null	float64
21	fractal_dimension_se	569 non-null	float64
22	radius_worst	569 non-null	float64
23	texture_worst	569 non-null	float64
24	perimeter_worst	569 non-null	float64
25	area_worst	569 non-null	float64
26	smoothness_worst	569 non-null	float64
27	compactness_worst	569 non-null	float64
28	concavity_worst	569 non-null	float64
29	concave points_worst	569 non-null	float64
30	symmetry_worst	569 non-null	float64
31	fractal_dimension_worst	569 non-null	float64
32	Unnamed: 32	0 non-null	float64

```
dtypes: float64(31), int64(1), object(1)
```

```
memory usage: 146.8+ KB
```

```
[3]: df.head()
```

```
[3]:      id diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  \
0   842302         M      17.99      10.38      122.80      1001.0
1   842517         M      20.57      17.77      132.90      1326.0
2  84300903         M      19.69      21.25      130.00      1203.0
3  84348301         M      11.42      20.38       77.58       386.1
4  84358402         M      20.29      14.34      135.10      1297.0

      smoothness_mean  compactness_mean  concavity_mean  concave points_mean  \
0         0.11840      0.27760      0.3001      0.14710
1         0.08474      0.07864      0.0869      0.07017
2         0.10960      0.15990      0.1974      0.12790
3         0.14250      0.28390      0.2414      0.10520
4         0.10030      0.13280      0.1980      0.10430

      ... texture_worst  perimeter_worst  area_worst  smoothness_worst  \
0   ...      17.33      184.60      2019.0      0.1622
1   ...      23.41      158.80      1956.0      0.1238
2   ...      25.53      152.50      1709.0      0.1444
3   ...      26.50       98.87       567.7      0.2098
4   ...      16.67      152.20      1575.0      0.1374

      compactness_worst  concavity_worst  concave points_worst  symmetry_worst  \
0         0.6656      0.7119      0.2654      0.4601
1         0.1866      0.2416      0.1860      0.2750
2         0.4245      0.4504      0.2430      0.3613
3         0.8663      0.6869      0.2575      0.6638
4         0.2050      0.4000      0.1625      0.2364

      fractal_dimension_worst  Unnamed: 32
0         0.11890      NaN
1         0.08902      NaN
2         0.08758      NaN
3         0.17300      NaN
4         0.07678      NaN
```

```
[5 rows x 33 columns]
```

1.1.3 For purposes of prediction, unique features such as the ID will not be useful and can be dropped

```
[4]: df1 = df.copy()
df1.drop('id',axis=1,inplace=True)
df1.head()
```

```
[4]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	M	17.99	10.38	122.80	1001.0	
1	M	20.57	17.77	132.90	1326.0	
2	M	19.69	21.25	130.00	1203.0	
3	M	11.42	20.38	77.58	386.1	
4	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

	symmetry_mean	...	texture_worst	perimeter_worst	area_worst	\
0	0.2419	...	17.33	184.60	2019.0	
1	0.1812	...	23.41	158.80	1956.0	
2	0.2069	...	25.53	152.50	1709.0	
3	0.2597	...	26.50	98.87	567.7	
4	0.1809	...	16.67	152.20	1575.0	

	smoothness_worst	compactness_worst	concavity_worst	concave points_worst	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	
2	0.1444	0.4245	0.4504	0.2430	
3	0.2098	0.8663	0.6869	0.2575	
4	0.1374	0.2050	0.4000	0.1625	

	symmetry_worst	fractal_dimension_worst	Unnamed: 32
0	0.4601	0.11890	NaN
1	0.2750	0.08902	NaN
2	0.3613	0.08758	NaN
3	0.6638	0.17300	NaN
4	0.2364	0.07678	NaN

[5 rows x 32 columns]

1.1.4 Since there aren't any missing values or categorical variables other than the diagnosis in the dataset, we can look at some basic descriptive statistics

```
[5]: df1.describe().T
```

```
[5]:
```

	count	mean	std	min	\
radius_mean	569.0	14.127292	3.524049	6.981000	
texture_mean	569.0	19.289649	4.301036	9.710000	
perimeter_mean	569.0	91.969033	24.298981	43.790000	
area_mean	569.0	654.889104	351.914129	143.500000	

smoothness_mean	569.0	0.096360	0.014064	0.052630
compactness_mean	569.0	0.104341	0.052813	0.019380
concavity_mean	569.0	0.088799	0.079720	0.000000
concave points_mean	569.0	0.048919	0.038803	0.000000
symmetry_mean	569.0	0.181162	0.027414	0.106000
fractal_dimension_mean	569.0	0.062798	0.007060	0.049960
radius_se	569.0	0.405172	0.277313	0.111500
texture_se	569.0	1.216853	0.551648	0.360200
perimeter_se	569.0	2.866059	2.021855	0.757000
area_se	569.0	40.337079	45.491006	6.802000
smoothness_se	569.0	0.007041	0.003003	0.001713
compactness_se	569.0	0.025478	0.017908	0.002252
concavity_se	569.0	0.031894	0.030186	0.000000
concave points_se	569.0	0.011796	0.006170	0.000000
symmetry_se	569.0	0.020542	0.008266	0.007882
fractal_dimension_se	569.0	0.003795	0.002646	0.000895
radius_worst	569.0	16.269190	4.833242	7.930000
texture_worst	569.0	25.677223	6.146258	12.020000
perimeter_worst	569.0	107.261213	33.602542	50.410000
area_worst	569.0	880.583128	569.356993	185.200000
smoothness_worst	569.0	0.132369	0.022832	0.071170
compactness_worst	569.0	0.254265	0.157336	0.027290
concavity_worst	569.0	0.272188	0.208624	0.000000
concave points_worst	569.0	0.114606	0.065732	0.000000
symmetry_worst	569.0	0.290076	0.061867	0.156500
fractal_dimension_worst	569.0	0.083946	0.018061	0.055040
Unnamed: 32	0.0	NaN	NaN	NaN

	25%	50%	75%	max
radius_mean	11.700000	13.370000	15.780000	28.11000
texture_mean	16.170000	18.840000	21.800000	39.28000
perimeter_mean	75.170000	86.240000	104.100000	188.50000
area_mean	420.300000	551.100000	782.700000	2501.00000
smoothness_mean	0.086370	0.095870	0.105300	0.16340
compactness_mean	0.064920	0.092630	0.130400	0.34540
concavity_mean	0.029560	0.061540	0.130700	0.42680
concave points_mean	0.020310	0.033500	0.074000	0.20120
symmetry_mean	0.161900	0.179200	0.195700	0.30400
fractal_dimension_mean	0.057700	0.061540	0.066120	0.09744
radius_se	0.232400	0.324200	0.478900	2.87300
texture_se	0.833900	1.108000	1.474000	4.88500
perimeter_se	1.606000	2.287000	3.357000	21.98000
area_se	17.850000	24.530000	45.190000	542.20000
smoothness_se	0.005169	0.006380	0.008146	0.03113
compactness_se	0.013080	0.020450	0.032450	0.13540
concavity_se	0.015090	0.025890	0.042050	0.39600
concave points_se	0.007638	0.010930	0.014710	0.05279

symmetry_se	0.015160	0.018730	0.023480	0.07895
fractal_dimension_se	0.002248	0.003187	0.004558	0.02984
radius_worst	13.010000	14.970000	18.790000	36.04000
texture_worst	21.080000	25.410000	29.720000	49.54000
perimeter_worst	84.110000	97.660000	125.400000	251.20000
area_worst	515.300000	686.500000	1084.000000	4254.00000
smoothness_worst	0.116600	0.131300	0.146000	0.22260
compactness_worst	0.147200	0.211900	0.339100	1.05800
concavity_worst	0.114500	0.226700	0.382900	1.25200
concave points_worst	0.064930	0.099930	0.161400	0.29100
symmetry_worst	0.250400	0.282200	0.317900	0.66380
fractal_dimension_worst	0.071460	0.080040	0.092080	0.20750
Unnamed: 32	NaN	NaN	NaN	NaN

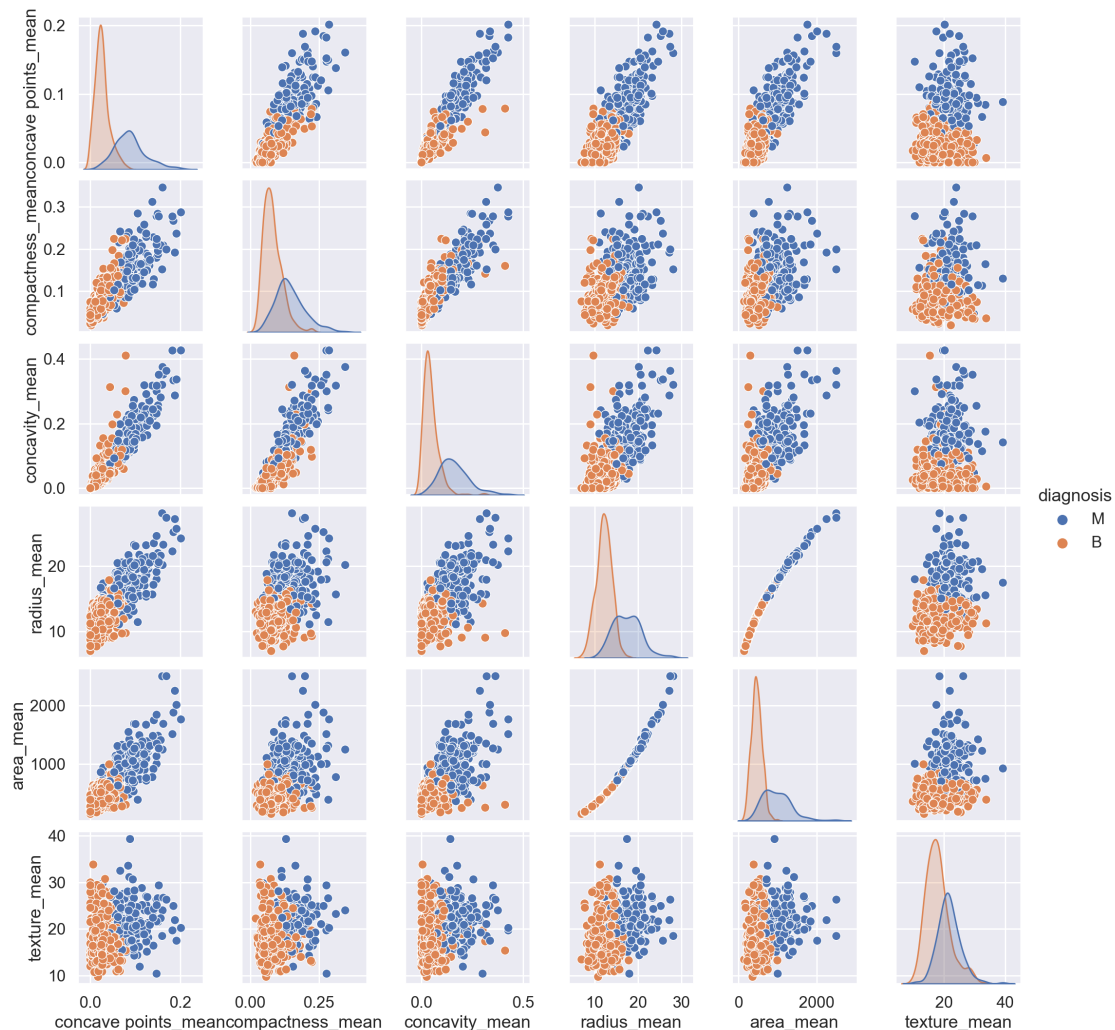
1.1.5 Whew, that's a lot of data. Maybe a pair plot with diagnosis as the hue can tell us more. Since the variables are so many, let's use only 6 of them for now.

```
[18]: smaller_df = df1.loc[:,['diagnosis','concave_
    ↪points_mean','compactness_mean','concavity_mean','radius_mean','area_mean','texture_mean']]
smaller_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   diagnosis              569 non-null    object
1   concave points_mean    569 non-null    float64
2   compactness_mean       569 non-null    float64
3   concavity_mean         569 non-null    float64
4   radius_mean            569 non-null    float64
5   area_mean              569 non-null    float64
6   texture_mean           569 non-null    float64
dtypes: float64(6), object(1)
memory usage: 31.2+ KB
```

```
[19]: sns.set_context('talk')
sns.pairplot(smaller_df, hue='diagnosis')
```

```
[19]: <seaborn.axisgrid.PairGrid at 0x18e204cab20>
```



1.1.6 There seems to be a linear relationship between dimensions and the diagnosis of the cancer. That would be an interesting hypothesis to test.

1.1.7 But we need to engineer these features into a normal distribution for linear regression modelling

```
[8]: # Check for skewed distributions before hypothesis testing
skew_vals = smaller_df.skew()
skew_limit = 0.75 # define a limit above which we will log transform
```

```
[9]: # Showing the skewed columns
skew_cols = (skew_vals
              .sort_values(ascending=False)
              .to_frame()
              .rename(columns={0: 'Skew'}))
```

```
.query('abs(Skew) > {}'.format(skew_limit)))
```

```
skew_cols
```

```
[9]:
```

	Skew
area_mean	1.645732
concavity_mean	1.401180
compactness_mean	1.190123
perimeter_mean	0.990650
radius_mean	0.942380

1.1.8 All features except texture are skewed. So we shall apply Log transformation to them before hypothesis testing

```
[10]: # Let's look at what happens to one of these features, when we apply np.log1p,
      ↪visually.

      # Choose a field
      field = "area_mean"

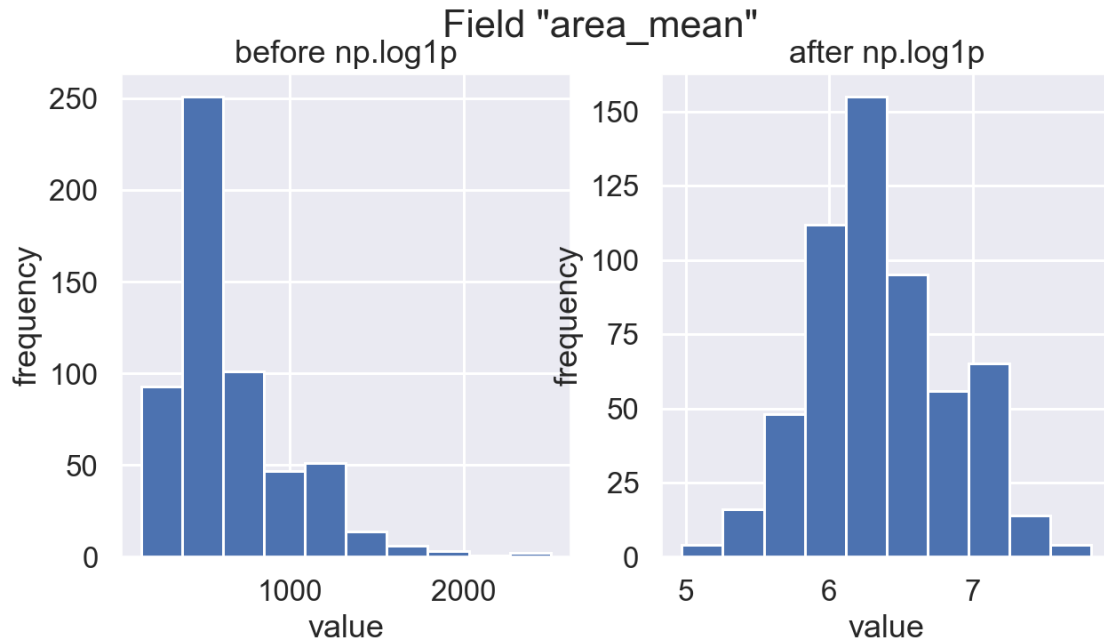
      # Create two "subplots" and a "figure" using matplotlib
      fig, (ax_before, ax_after) = plt.subplots(1, 2, figsize=(10, 5))

      # Create a histogram on the "ax_before" subplot
      df[field].hist(ax=ax_before)

      # Apply a log transformation (numpy syntax) to this column
      df[field].apply(np.log1p).hist(ax=ax_after)

      # Formatting of titles etc. for each subplot
      ax_before.set(title='before np.log1p', ylabel='frequency', xlabel='value')
      ax_after.set(title='after np.log1p', ylabel='frequency', xlabel='value')
      fig.suptitle('Field "{}".format(field))
```

```
[10]: Text(0.5, 0.98, 'Field "area_mean"')
```

```
[11]: # Perform the skew transformation:
```

```
for col in skew_cols.index.values:
    if col == "diagnosis":
        continue
    smaller_df[col] = smaller_df[col].apply(np.log1p)
```

```
[12]: smaller_df.head()
```

```
[12]:
```

	diagnosis	perimeter_mean	compactness_mean	concavity_mean	radius_mean	\
0	M	4.818667	0.244983	0.262441	2.943913	
1	M	4.897093	0.075701	0.083330	3.071303	
2	M	4.875197	0.148334	0.180153	3.029650	
3	M	4.364117	0.249902	0.216240	2.519308	
4	M	4.913390	0.124692	0.180653	3.058237	

	area_mean	texture_mean
0	6.909753	10.38
1	7.190676	17.77
2	7.093405	21.25
3	5.958683	20.38
4	7.168580	14.34

```
[13]: smaller_df.describe().T
```

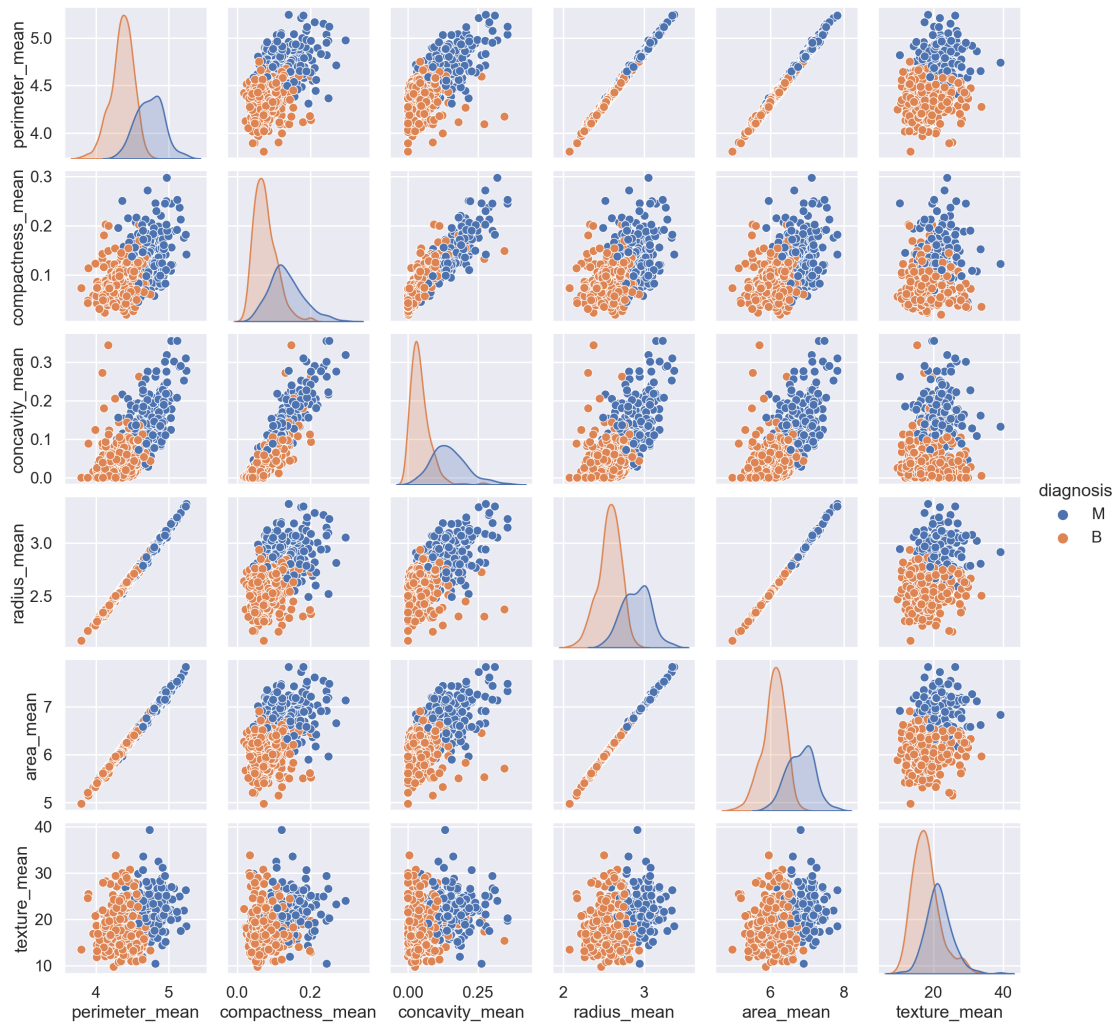
```
[13]:
```

	count	mean	std	min	25%	50% \
perimeter_mean	569.0	4.500683	0.248328	3.801985	4.332968	4.468663
compactness_mean	569.0	0.098145	0.046646	0.019195	0.062900	0.088588
concavity_mean	569.0	0.082552	0.070113	0.000000	0.029132	0.059721
radius_mean	569.0	2.691235	0.222226	2.077064	2.541602	2.665143
area_mean	569.0	6.365109	0.482274	4.973280	6.043345	6.313729
texture_mean	569.0	19.289649	4.301036	9.710000	16.170000	18.840000

	75%	max
perimeter_mean	4.654912	5.244389
compactness_mean	0.122572	0.296691
concavity_mean	0.122837	0.355434
radius_mean	2.820188	3.371082
area_mean	6.664026	7.824846
texture_mean	21.800000	39.280000

```
[14]: # Checking to see how the feature engineering has improved the distribution of
      ↪ attributes
      sns.pairplot(smaller_df, hue='diagnosis')
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x18e205b9730>
```



1.2 Hypotheses

Ho - The concave points mean of benign and malignant tissue is the same

Ho - The compactness mean of benign and malignant tissue is the same

Ho - The perimeter mean of benign and malignant tissue is the same

1.2.1 Hypothesis testing

Ho - The perimeter mean of benign and malignant tissue is the same

H1 - The perimeter mean of benign and malignant tissue is not the same

[15]: *# To test this hypothesis, we shall carry out a t-test on the two diagnoses*

```
from scipy.stats import ttest_ind
```

```
mal = smaller_df[smaller_df['diagnosis'] == 'M']
ben = smaller_df[smaller_df['diagnosis'] == 'B']

ttest_ind(mal['perimeter_mean'],ben['perimeter_mean'])
```

```
[15]: Ttest_indResult(statistic=26.35462190815133, pvalue=1.5372483643273375e-100)
```

Such a low p-value leads us to reject the null hypothesis. Tissues of benign and malignant breast cancer have different perimeters. Next steps;

1. Check how other features correlate with the cancer being benign or malignant.
2. This data set contained only 569 samples. More data would be needed to confirm or refute these findings
3. A model can be trained to look for these key features in identifying benign or malignant cancer tissue.

Thank you!